



UNIVERSITY OF TEHRAN

COLLEGE OF ENGINEERING

DEPARTMENT OF ELECTRICAL AND COMPUTER ENGINEERING

NEURAL NETWORK & DEEP LEARNING
STOCK MARKET PREDICTION USING RNN, LSTM, GRU

SIAVASH SHAMS

MOHAMMAD HEYDARI

SCHOOL OF ELECTRICAL AND COMPUTER ENGINEERING

UNIVERSITY OF TEHRAN

May. 2022

1 **CONTENTS**

Stock Market Prediction3

 1.1 3



 1.2 11

 1.3 18

STOCK MARKET PREDICTION

1.1

Below we can see Google and Apple, day to day stock closing prices from 2010 to 2018.




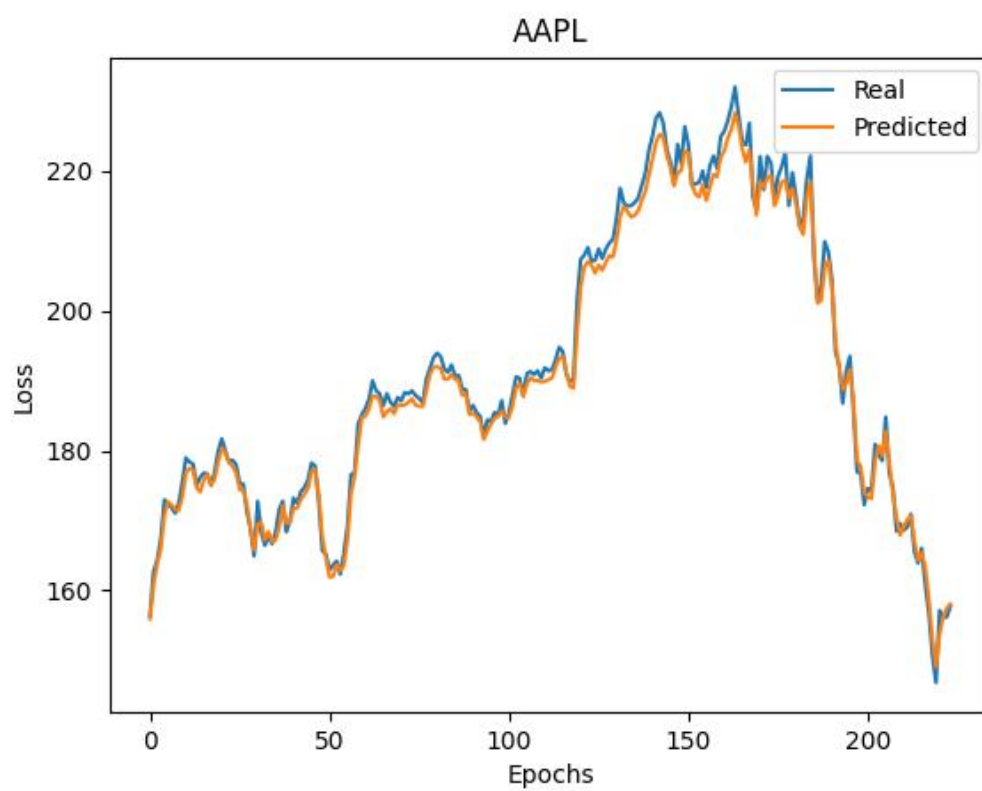

- **LSTM**

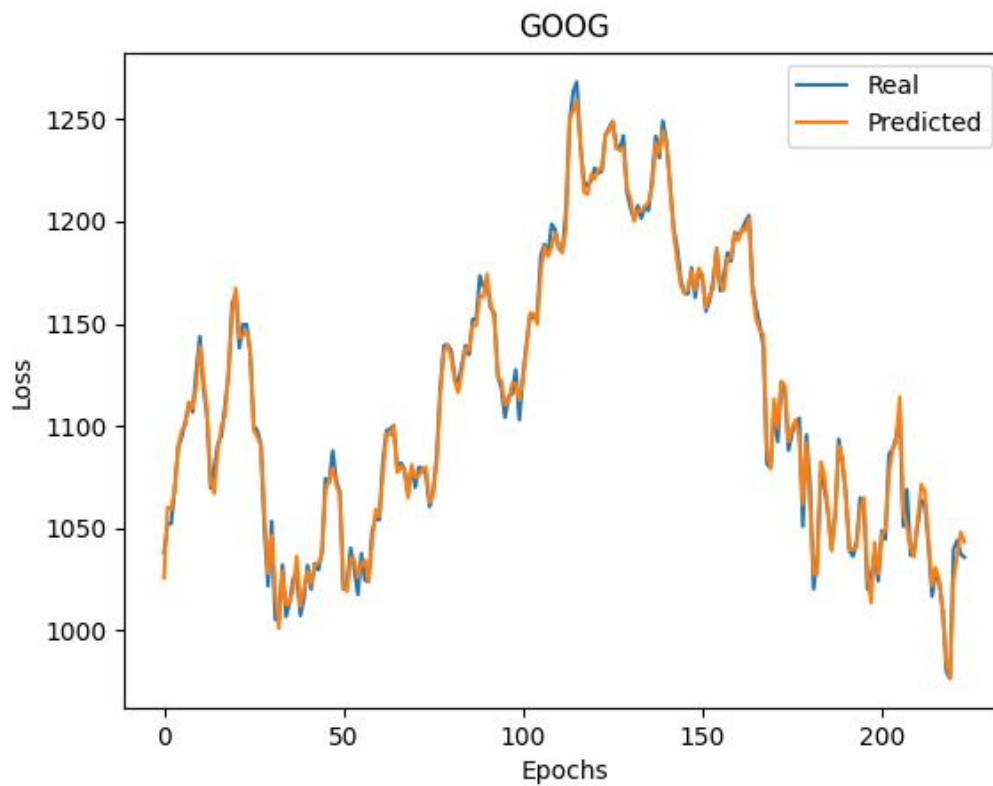
```
Model: "sequential"
-----
Layer (type)                 Output Shape          Param #
-----
lstm (LSTM)                   (None, 30, 64)        19712
-----
lstm_1 (LSTM)                 (None, 64)            33024
-----
dense (Dense)                 (None, 2)             130
-----
Total params: 52,866
Trainable params: 52,866
Non-trainable params: 0
```

We will train our model using Adam optimizer with 100 epochs.

```
Total training time: 106.46464157104492 seconds
```







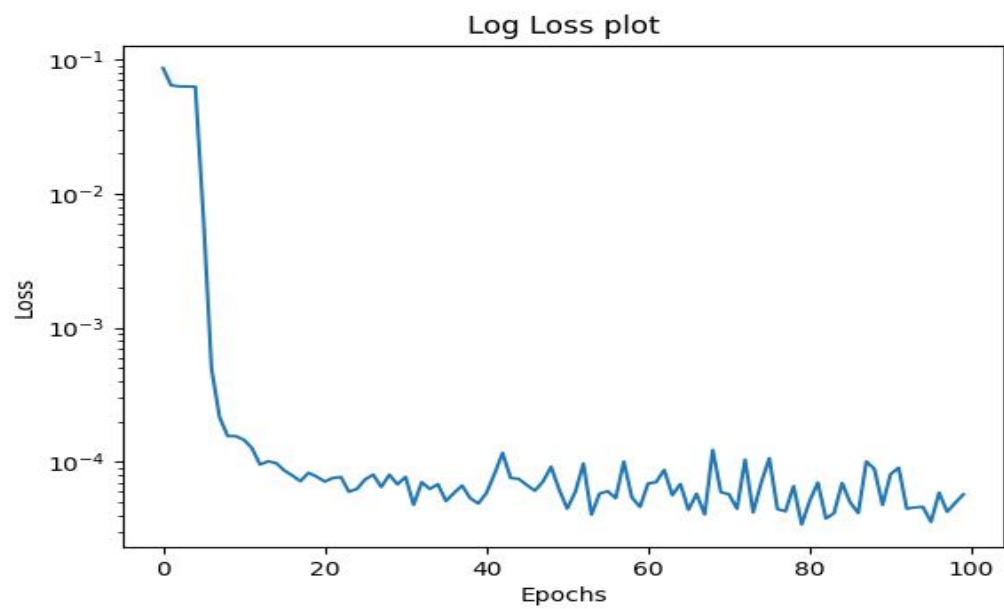

- RNN

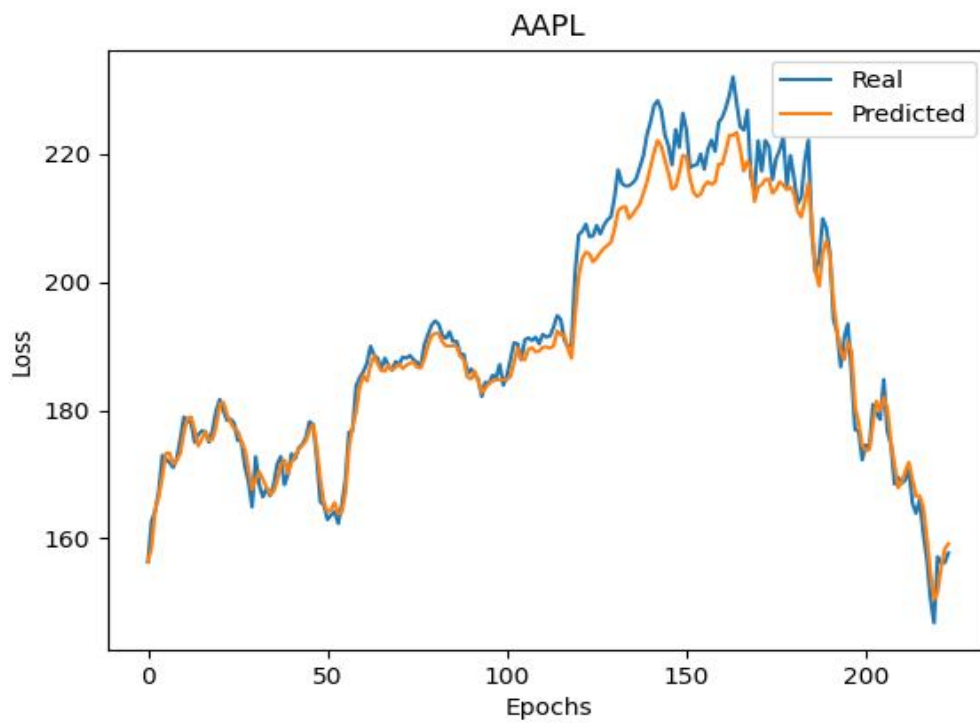

Model: "sequential_1"

Layer (type)	Output Shape	Param #
simple_rnn (SimpleRNN)	(None, 30, 64)	4928
simple_rnn_1 (SimpleRNN)	(None, 64)	8256
dense_1 (Dense)	(None, 2)	130

=====
 Total params: 13,314
 Trainable params: 13,314
 Non-trainable params: 0

```
Total training time: 26.431206941604614 seconds
```






- GRU

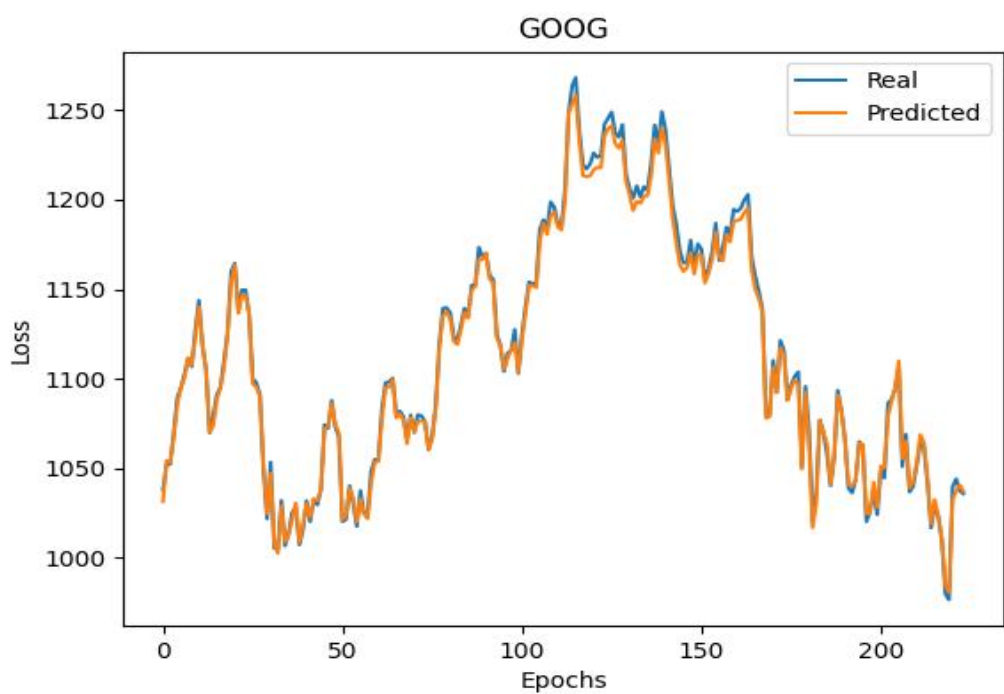

```
Model: "sequential_1"
```


Layer (type)	Output Shape	Param #
gru (GRU)	(None, 30, 64)	14976
gru_1 (GRU)	(None, 64)	24960
dense_1 (Dense)	(None, 2)	130

Total params: 40,066
Trainable params: 40,066
Non-trainable params: 0

```
Total training time: 81.92196726799011 seconds
```







By looking at the results we conclude that GRU outperforms the other two memory cells

Performance: $GRU > LSTM > RNN$



also training time is least for RNN and GRU comes second, and LSTM takes the longest. We can see that as the memory unit complexity (number of parameters) gets bigger, it requires more training time.



Training time: $LSTM > GRU > RNN$

1.2

- GRU + MAPE

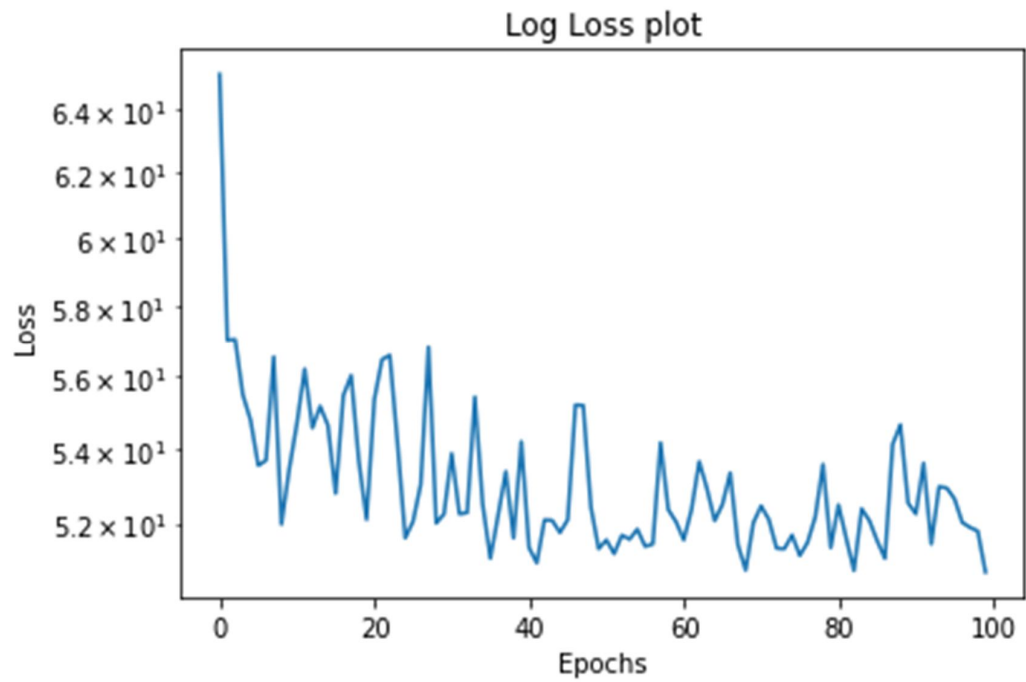

```
Total training time: 145.55148577690125 seconds
```

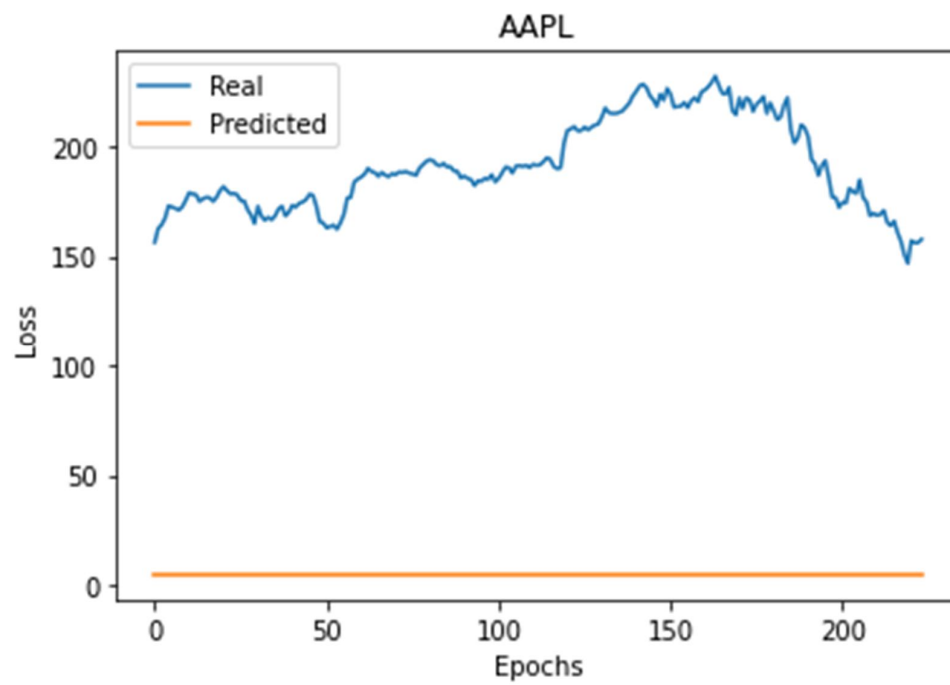
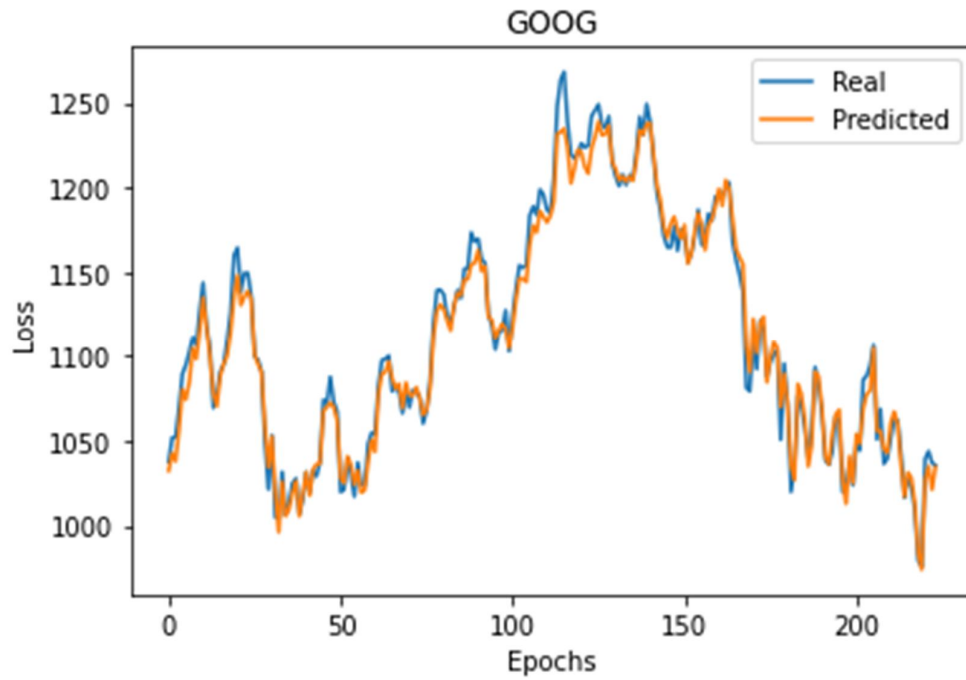




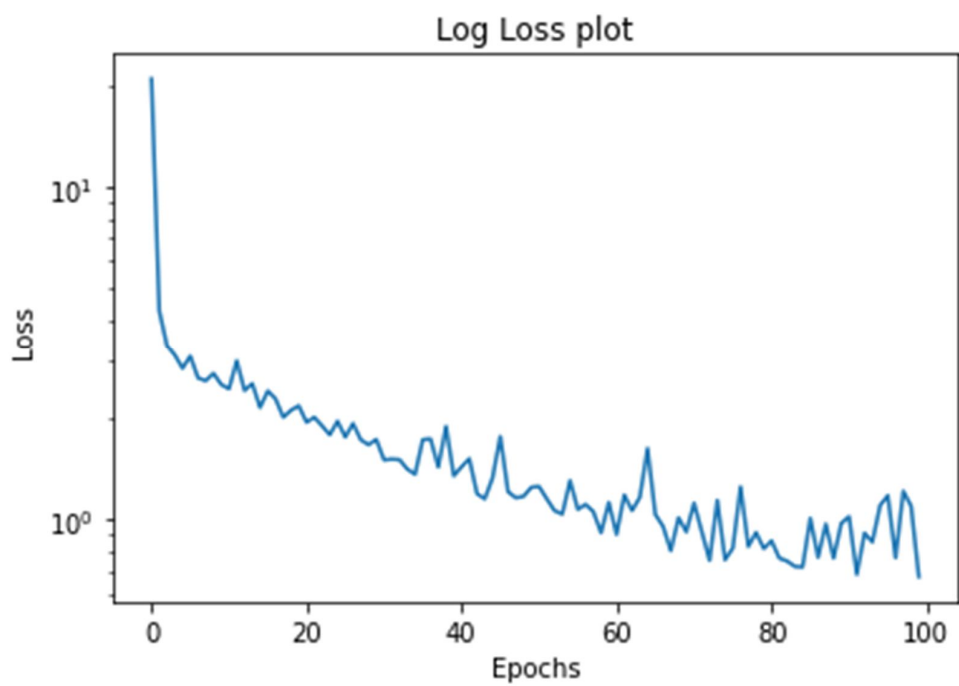

- MAPE + RNN

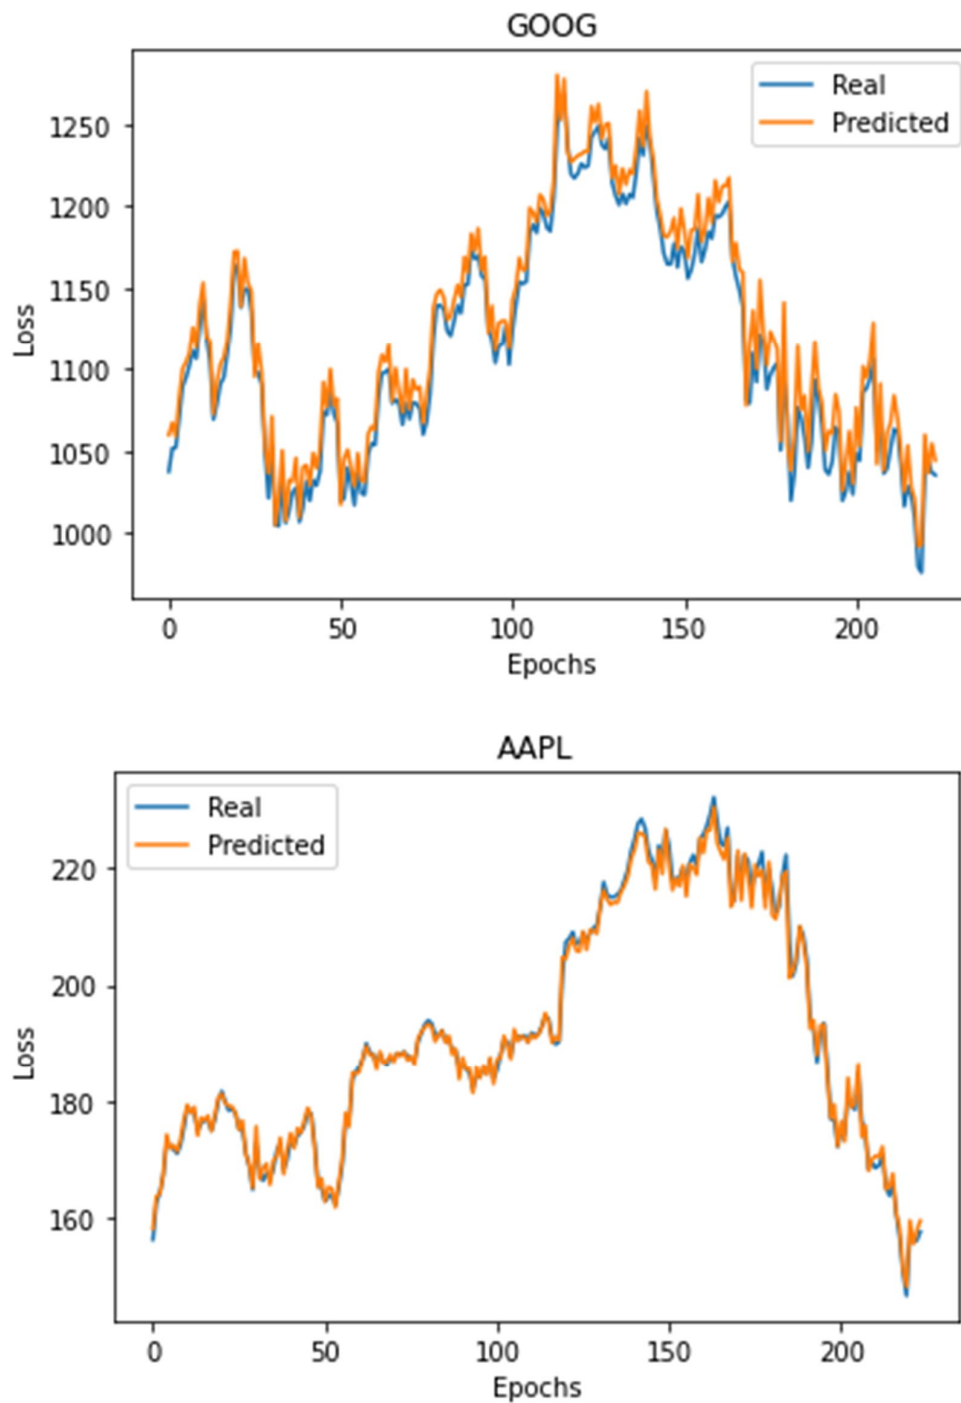
```
Total training time: 60.38407301902771 seconds
```





- MAPE + LSTM







As we can see from the results, the MAPE loss function doesn't improve GRU and LSTM but reduces performance of RNN, as it has a constant gradient it suffers from vanishing gradient and thus the model is not trained well for RNN

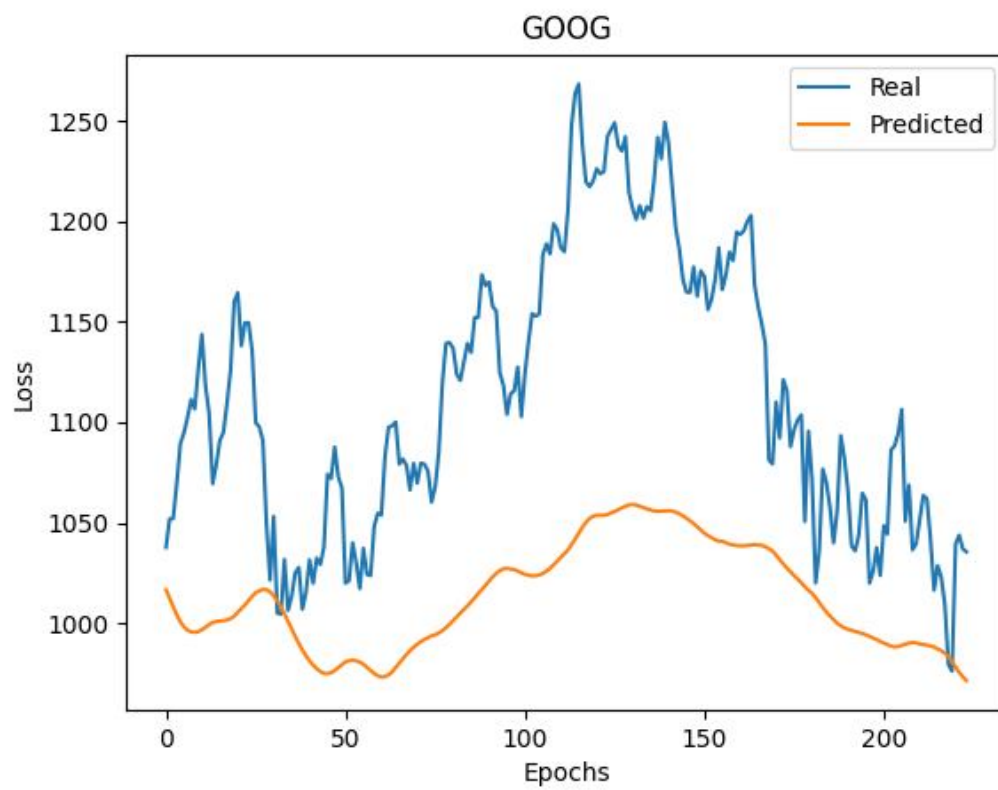
1.3


LSTM + Adagrad

```
Total training time: 149.08284449577332 seconds
```













- LSTM + RMSprop

```
Total training time: 205.3320653438568 seconds
```











- ADAGRAD + RNN



```
Total training time: 59.3671555519104 seconds
```








- RMSprop +RNN



```
Total training time: 63.164148807525635 seconds
```







- ADAgrad + GRU



```
Total training time: 137.29109835624695 seconds
```






- RMSprop + GRU

```
Total training time: 137.48347330093384 seconds
```







As we can see from the results both ADAGRAD and RMSprop optimizers get outperformed by Adam, because Adam uses both first and second moments and benefits from advantages of two other optimizers, so it is generally the best choice.


Also, Adam takes noticeably less time to train.


Optimiser	Year	Learning Rate	Gradient
Momentum	1964		✓
AdaGrad	2011	✓	
RMSprop	2012	✓	
Adadelta	2012	✓	
Nesterov	2013		✓
Adam	2014	✓	✓
AdaMax	2015	✓	✓
Nadam	2015	✓	✓
AMSGrad	2018	✓	✓

- LSTM + 0.1 Dropout













RNN + dropout 0.1













GRU + 0.1 dropout









By comparing the results to part 2 we can see that adding dropout decreased performance of LSTM and GRU units but improved RNN unit.

Dropout is a regularization technique, and is most effective at preventing overfitting. However, there are several places when dropout can hurt performance.

- when the network is small relative to the dataset, regularization is usually unnecessary. If the model capacity is already low, lowering it further by adding regularization will hurt performance.
- when training time is limited. Usually, dropout hurts performance at the start of training, but results in the final "converged" error being lower. Therefore, if you don't plan to train until convergence, you may not want to use dropout.